



Computer Vision and Image Processing (CSEL-393)

Lecture


Dr. Qurat ul Ain Akram
 Assistant Professor
 Computer Science Department (New
 Campus) KSK, UET, Lahore

Edge detection



- Goal: Identify sudden changes (discontinuities) in an image
 - Intuitively, most semantic and shape information from the image can be encoded in the edges
 - More compact than pixels
- Ideal: artist's line drawing (but artist is also using object-level knowledge)

Edge Detection



- Convert a 2D image into a set of curves
 - Extracts salient features of the scene
 - More compact than pixels


Application

- What is an object
- How can we find it




Edge Detection in images

- At edges intensity or color changes




Origin of Edges




- Edges are caused by a variety of factors


Evaluation Metrics




Evaluation Metrics




Evaluation Metrics




Evaluation Metrics




Evaluation Metrics




Evaluation Metrics




Edge Types



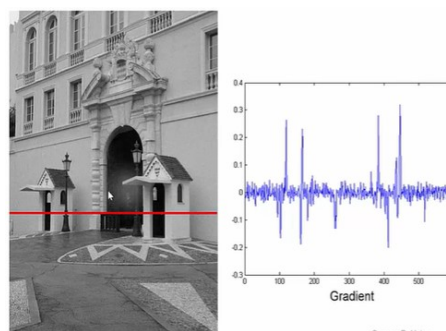
Example of Edges



- An edge is a place of rapid change in the image intensity function




Effect of noise



Effect of Noise

- Consider a single row or column of the image
- Plotting intensity as a function of position gives a signal
- Where is the edge




Effects of Noise

- Difference filters respond strongly to noise
 - Image noise results in pixels that look very different from their neighbors
 - Generally, the larger the noise the stronger the response
- What can we do about it?

Solution

- First Smooth the image




- To find edges, look for peaks in $\frac{d}{dx}(f * g)$

Derivative Theorem of Smoothing


- Differentiation is convolution, and convolution is associative:

$$\frac{d}{dx}(f * g) = f * \frac{d}{dx}g$$

- This saves us one operation:



Tradeoff between smoothing and localization




- Smoothed derivative removes noise, but blurs edge. Also finds edges at different "scales".

Derivatives and Noise

- Strongly affected by noise**
 - obvious reason: image noise results in pixels that look very different from their neighbors
- The larger the noise is the stronger the response
- What is to be done?**
 - Neighboring pixels look alike
 - Pixel along an edge look alike
 - Image smoothing should help
 - Force pixels different from their neighbors (possibly noise) to look like neighbors

Derivatives and Noise



Edge Detectors


- Gradient Operator
 - Prewitt
 - Sobel
- Laplacian of Gaussian
- Gradient of Gaussian (Canny Edge Detector)

Prewitt and Sobel Edge Detector

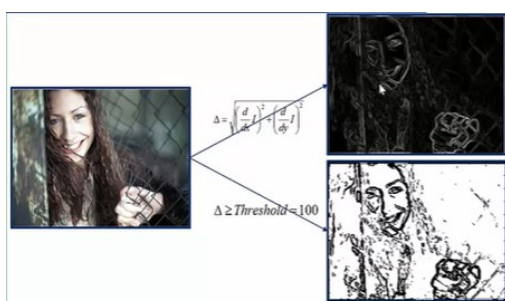


- Compute derivatives in x and y directions
- Find gradient magnitude
- Threshold gradient magnitude



Prewitt Edge Detector



Sobel Edge Detector




Marr Hildreth Edge Detector

- Smooth image by Gaussian filter $\rightarrow S$
- Apply Laplacian to S
 - Used in mechanics, electromagnetics, wave theory, quantum mechanics and Laplace equation
- Find zero crossings
 - Scan along each row, record an edge point at the location of zero-crossing
 - Repeat above step along each column

Marr Hildreth Edge Detector

- Gaussian smoothing

$$\hat{S} = \hat{g} * \hat{I} \quad g = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2+y^2}{2\sigma^2}}$$



- Find Laplacian

$$\Delta^2 S = \frac{\partial^2}{\partial x^2} S + \frac{\partial^2}{\partial y^2} S$$

second order derivative in x second order derivative in y
 • ∇ is used for gradient (first derivative)
 • Δ^2 is used for Laplacian (Second derivative)

Laplacian of Gaussian

- Deriving the Laplacian of Gaussian (LoG)

$$\Delta^2 S = \Delta^2 (g * I) = (\Delta^2 g) * I \quad g = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2+y^2}{2\sigma^2}}$$


$$g_x = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{x^2+y^2}{2\sigma^2}} \left(-\frac{2x}{2\sigma^2} \right)$$

$$\Delta^2 g = -\frac{1}{\sqrt{2\pi}\sigma^3} \left(2 - \frac{x^2+y^2}{\sigma^2} \right) e^{-\frac{x^2+y^2}{2\sigma^2}}$$

Finding Zero Crossing


- Four cases of zero-crossings :
 - {+, -}
 - {+, 0, -}
 - {-, +}
 - {-, 0, +}
- Slope of zero-crossing {a, -b} is |a+b|.
- To mark an edge
 - compute slope of zero-crossing
 - Apply a threshold to slope

Example



Example

$$\Delta^2 G_\sigma = -\frac{1}{\sqrt{2\pi}\sigma^3} \left(2 - \frac{x^2 + y^2}{\sigma^2} \right) e^{-\frac{x^2 + y^2}{2\sigma^2}}$$



LOG Algorithm

- Apply LOG to the image
- Find Zero crossings of each row
- Find slope of zero crossing
- Apply threshold to the slope and mark edges

Canny Edge Detection

- Canny Edge Detector Steps
 1. Smooth image with Gaussian filter
 2. Compute derivative of filtered image
 3. Find magnitude and orientation of gradient
 4. Apply "Non-maximum Suppression"
 5. Apply "Hysteresis Threshold" (use range between low and high)

Home assignment


- Write a python code
 - Read an image
 - Find edges using
 1. Prewitt and sobel
 2. Laplacian of Gaussian (LOG)
 3. Canny
- Write image having marked edges on drive

Readings

- Chapter
- Richard Szeliski, Computer Vision, Algorithms and Applications, 2nd Ed, <https://szeliski.org/Book/>

Gradient

- Gradient equation: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$
- Represents direction of most rapid change in intensity



- Gradient direction: $\theta = \tan^{-1} \left(\frac{\partial f / \partial y}{\partial f / \partial x} \right)$
- The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Discrete Edge Operators

- How can we differentiate a **discrete** image?

Finite difference approximations:

$$\frac{\partial I}{\partial x} \approx \frac{1}{2\varepsilon} ((I_{i+1,j+1} - I_{i,j+1}) + (I_{i+1,j} - I_{i,j}))$$

$$\frac{\partial I}{\partial y} \approx \frac{1}{2\varepsilon} ((I_{i+1,j+1} - I_{i+1,j}) + (I_{i,j+1} - I_{i,j}))$$

Convolution masks :

$$\frac{\partial I}{\partial x} \approx \frac{1}{2\varepsilon} \begin{bmatrix} -1 & 1 \\ -1 & 1 \end{bmatrix} \quad \frac{\partial I}{\partial y} \approx \frac{1}{2\varepsilon} \begin{bmatrix} 1 & 1 \\ -1 & -1 \end{bmatrix}$$

Discrete Edge Operators

- First order partial derivatives:

$$\frac{\partial I}{\partial x} \approx \begin{bmatrix} -1 & 1 \\ -1 & 1 \end{bmatrix} \quad \frac{\partial I}{\partial y} \approx \begin{bmatrix} 1 & 1 \\ -1 & -1 \end{bmatrix}$$

- Second order partial derivatives:

- Laplacian :

$$\nabla^2 I = \frac{\partial^2 I}{\partial x^2} + \frac{\partial^2 I}{\partial y^2}$$

Convolution masks :

$$\nabla^2 I \approx \begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix} \quad \text{or} \quad \begin{bmatrix} 1 & 4 & 1 \\ 4 & -20 & 4 \\ 1 & 4 & 1 \end{bmatrix} \quad (\text{more accurate})$$

The Sobel Operators

- Better approximations of the gradients exist
 - The *Sobel* operators below are commonly used

$$s_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$s_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$$

Comparing Edge Operators

Gradient: $\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$ Good Localization
Noise Sensitive
Poor Detection

Roberts (2 x 2):

$$\begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix} \quad \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$$

Sobel (3 x 3):

$$\begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$



Sobel (5 x 5):

$$\begin{bmatrix} -1 & -2 & 0 & 2 & 1 \\ -2 & -3 & 0 & 3 & 2 \\ -3 & -5 & 0 & 5 & 3 \\ -2 & -3 & 0 & 3 & 2 \\ -1 & -2 & 0 & 2 & 1 \end{bmatrix} \quad \begin{bmatrix} 1 & 2 & 3 & 2 & 1 \\ 2 & 3 & 5 & 3 & 2 \\ 0 & 0 & 0 & 0 & 0 \\ -2 & -3 & -5 & -3 & -2 \\ -1 & -2 & -3 & -2 & -1 \end{bmatrix}$$

Poor Localization
Less Noise Sensitive
Good Detection


Effects of Noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Where is the edge??

Solution: Smooth First




Where is the edge?

Look for peaks in $\frac{\partial}{\partial x}(h * f)$


Derivative Theorem of Convolution

$$\frac{\partial}{\partial x}(h * f) = \left(\frac{\partial}{\partial x}h\right) * f \quad \dots \text{saves us one operation.}$$



Laplacian of Gaussian (LoG)

$$\frac{\partial^2}{\partial x^2}(h * f) = \left(\frac{\partial^2}{\partial x^2}h\right) * f \quad \text{Laplacian of Gaussian}$$



Where is the edge?

Zero-crossings of bottom graph !

Canny Edge Operator

- Smooth image I with 2D Gaussian: $G * I$
- Find local edge normal directions for each pixel

$$\bar{n} = \frac{\nabla(G * I)}{|\nabla(G * I)|}$$

- Compute edge magnitudes $|\nabla(G * I)|$
- Locate edges by finding zero-crossings along the edge normal directions (**non-maximum suppression**)

$$\frac{\partial^2(G * I)}{\partial \bar{n}^2} = 0$$

The Canny Edge Detector



original image (Lena)

The Canny Edge Detector



magnitude of the gradient

The Canny Edge Detector



After non-maximum suppression

Canny Edge Operator



original

Canny with $\sigma = 1$

Canny with $\sigma = 2$

- The choice of σ depends on desired behavior
 - large σ detects large scale edges
 - small σ detects fine features